

Introduction to Machine Learning (Spring 2019)

Homework #4 (50 Pts, May 22)

Student ID _____

Name _____

Instruction: We provide all codes and datasets in Python. Please write your code to complete Perceptron & MLP. **Compress ‘Answer.py’ & your report ONLY and submit with the filename ‘HW2_STUDENT_ID.zip’.**

(1) [30 pts] Implement Perceptron & MLP in ‘Answer.py’.

(a)[Perceptron, 10 pts] Implement sign function and perceptron in ‘Answer.py’ (‘sign’, ‘Perceptron’).

```
def sign(z):
    sign_z = None
    # ===== EDIT HERE =====
    save = z.shape
    sign_z = z.reshape(-1)

    for ind in range(len(sign_z)):
        sign_z[ind] = 1 if sign_z[ind] > 0 else -1

    sign_z = sign_z.reshape(save)
    # =====
    return sign_z

class Perceptron:
    def __init__(self, num_features):
        # NOTE : In this assignment, weight and bias are separated. Be careful.
        self.W = np.random.rand(num_features, 1)
        self.b = np.random.rand(1)

    def forward(self, x):
        out = None
        if len(x.shape) < 2:
            x = np.expand_dims(x, 0)
        # ===== EDIT HERE =====
        num_data = x.shape[0]
        out = sign(np.matmul(x, self.W) + self.b)
        # =====
        return out
```

```

def stochastic_train(self, x, y, learning_rate):
    num_data = x.shape[0]
    while True:
        # Repeat until quit condition is satisfied.
        quit = True

        for i in range(num_data):
            # ===== EDIT HERE =====
            out = self.forward(x[i])
            if out != y[i]:
                self.W += np.multiply(x[i], y[i] * learning_rate)
                self.b += y[i] * learning_rate
                quit = False

            # =====

        if quit:
            break

def batch_train(self, x, y, learning_rate):
    num_data = x.shape[0]
    while True:
        # gradients of W & b
        dW = np.zeros_like(self.W)
        db = np.zeros_like(self.b)

        # Repeat until quit condition is satisfied.
        quit = True
        for i in range(num_data):
            # ===== EDIT HERE =====
            out = self.forward(x[i])
            if out != y[i]:
                dW += np.reshape(x[i] * y[i] * learning_rate, (-1, 1))
                db += y[i] * learning_rate
                quit = False

        self.W += dW
        self.b += db
        # =====

        if quit:
            break

```

(b) [MLP, 20 pts] Implement activation functions and MLP layers in ‘Answer.py’ (‘Sigmoid’, ‘ReLU’, ‘Input/Hidden/(Sigmoid, Softmax) Output Layers’).

class ReLU:

```
def __init__(self):
    # 1 (True) if ReLU input < 0
    self.zero_mask = None

def forward(self, z):
    out = None
    # ===== EDIT HERE =====
    save = z.shape
    temp = z.reshape(-1)
    mask = np.zeros_like(temp)

    for ind in range(len(temp)):
        mask[ind] = 1 if temp[ind] >= 0 else 0
        temp[ind] = temp[ind] if temp[ind] >= 0 else 0

    out = temp.reshape(save)
    self.zero_mask = mask.reshape(save)
    # =====
    return out

def backward(self, d_prev):
    dz = None
    # ===== EDIT HERE =====

    dz = np.multiply(d_prev, self.zero_mask)
    # =====
    return dz
```

class Sigmoid:

```
def __init__(self):
    self.out = None

def forward(self, z):
    self.out = None
    # ===== EDIT HERE =====
    save = z.shape
    temp = z.reshape(-1)

    for ind in range(len(temp)):
        temp[ind] = 1 / (1 + np.exp(-temp[ind]))

    self.out = temp.reshape(save)
    # =====
    return self.out

def backward(self, d_prev):
    dz = None
    # ===== EDIT HERE =====
    dz = np.multiply(self.out, np.add(1, -self.out))
    dz = np.multiply(dz, d_prev)
    # =====
    return dz
```

```

class InputLayer:
    """
    def __init__(self, num_features, num_hidden_1, activation):
        # Weights and bias
        self.W = np.random.rand(num_features, num_hidden_1)
        self.b = np.zeros(num_hidden_1)
        # Gradient of Weights and bias
        self.dW = None
        self.db = None
        # Forward input
        self.x = None
        # Activation function (Sigmoid or ReLU)
        self.act = activation()

    def forward(self, x):
        self.x = None
        self.out = None
        # ===== EDIT HERE =====
        self.x = x
        inner = np.add(np.matmul(self.x, self.W), self.b)

        self.out = self.act.forward(inner)
        # =====
        return self.out

    def backward(self, d_prev):
        self.dW = None
        self.db = None
        # ===== EDIT HERE =====
        act_backward = self.act.backward(d_prev)

        self.dW = np.matmul(np.transpose(self.x), act_backward)
        self.db = np.sum(act_backward, axis=0)
        # =====

```

```

class SigmoidOutputLayer:
    def __init__(self, num_hidden_2, num_outputs):
        # Weights and bias
        self.W = np.random.rand(num_hidden_2, num_outputs)
        self.b = np.zeros(num_outputs)
        # Gradient of Weights and bias
        self.dW = None
        self.db = None
        # Input (x), label(y), prediction(y_hat)
        self.x = None
        self.y = None
        self.y_hat = None
        # Loss
        self.loss = None
        # Sigmoid function
        self.sigmoid = Sigmoid()

    def forward(self, x, y):
        self.y_hat = self.predict(x)
        self.y = y
        self.x = x
        self.loss = self.binary_ce_loss(self.y_hat, self.y)
        return self.loss

```

```

def binary_ce_loss(self, y_hat, y):
    eps = 1e-10
    bce_loss = None
    # ===== EDIT HERE =====
    batch_size = y.shape[0]

    temp_y = y.reshape(-1)
    temp_y_hat = y_hat.reshape(-1)

    bce_loss = 0
    for i in range(len(temp_y)):
        bce_loss -= temp_y[i] * np.log(temp_y_hat[i] + eps) +W
            (1 - temp_y[i]) * np.log(1 - temp_y_hat[i] + eps)

    bce_loss /= batch_size
    # =====
    return bce_loss

def predict(self, x):
    y_hat = None
    # ===== EDIT HERE =====
    z = np.matmul(x, self.W) + self.b
    y_hat = self.sigmoid.forward(z)
    # =====
    return y_hat

def backward(self, d_prev=1):
    batch_size = self.y.shape[0]
    dx = None
    # ===== EDIT HERE =====

    # This equation is derived from the derivative of cross entropy w.r.t y_hat
    # Letting y_hat = c, the gradient is (c - y) / ( (c * (1 - c)) * batch_size )

    y_diff = np.multiply(self.y_hat, 1 - self.y_hat)
    loss_grad = np.divide((self.y_hat - self.y), y_diff * batch_size)
    sig_backward = self.sigmoid.backward(d_prev * loss_grad)

    self.dW = np.matmul(np.transpose(self.x), sig_backward)
    self.db = np.sum(sig_backward, axis=0)

    dx = np.transpose(np.matmul(self.W, np.transpose(sig_backward)))

    # =====

    return dx

```

```

class HiddenLayer:
    def __init__(self, num_hidden_1, num_hidden_2):
        # Weights and bias
        self.W = np.random.rand(num_hidden_1, num_hidden_2)
        self.b = np.zeros(num_hidden_2)
        # Gradient of Weights and bias
        self.dW = None
        self.db = None
        # ReLU function
        self.act = ReLU()

    def forward(self, x):
        self.x = None
        self.out = None
        # ===== EDIT HERE =====
        self.x = x
        inner = np.add(np.matmul(self.x, self.W), self.b)

        self.out = self.act.forward(inner)
        # =====
        return self.out

    def backward(self, d_prev):
        dx = None
        self.dW = None
        self.db = None
        # ===== EDIT HERE =====
        act_backward = self.act.backward(d_prev)

        self.dW = np.matmul(np.transpose(self.x), act_backward)
        self.db = np.sum(act_backward, axis=0)

        dx = np.transpose(np.matmul(self.W, np.transpose(act_backward)))

        # =====
        return dx

class SoftmaxOutputLayer:
    def __init__(self, num_hidden_2, num_outputs):
        # Weights and bias
        self.W = np.random.rand(num_hidden_2, num_outputs)
        self.b = np.zeros(num_outputs)
        # Gradient of Weights and bias
        self.dW = None
        self.db = None
        # Input (x), label(y), prediction(y_hat)
        self.x = None
        self.y = None
        self.y_hat = None
        # Loss
        self.loss = None

```

```
def forward(self, x, y):
    self.y_hat = self.predict(x)
    self.y = y
    self.x = x

    self.loss = self.ce_loss(self.y_hat, self.y)

    return self.loss
```

```
def ce_loss(self, y_hat, y):
    eps = 1e-10
    ce_loss = None
    # ===== EDIT HERE =====
    batch_size = y.shape[0]

    pred_log = np.log(y_hat + eps)
    ce_loss = np.sum(np.multiply(y, -pred_log))
    ce_loss /= batch_size
    # =====
    return ce_loss
```

```
def predict(self, x):
    y_hat = None
    # ===== EDIT HERE =====

    z = np.matmul(x, self.W) + self.b
    y_hat = softmax(z)

    # =====
    return y_hat
```

```
def backward(self, d_prev=1):
    batch_size = self.y.shape[0]
    dx = None
    # ===== EDIT HERE =====

    # Sim as sigmoid
    loss_grad = np.divide(self.y_hat - self.y, batch_size)

    self.dW = np.matmul(np.transpose(self.x), loss_grad)
    self.db = np.sum(loss_grad, axis=0)

    dx = np.transpose(np.matmul(self.W, np.transpose(loss_grad)))

    # =====
    return dx
```

```
# =====
```

(2) [20 Pts] Experiment results

- (a) [MLP-1] Adjust 'num_epochs' and 'learning_rate' and run 'MLP_1.py' to solve XOR problem. Report training accuracy with given code and explain how the MLP solve XOR problem by analyzing values of hidden nodes.


```
# ===== EDIT HERE =====  
num_epochs = 1000  
learning_rate = 1  
print_every = True  
# =====
```

Result :

```
==== [TEST] ====  
Pred : 0, Answer 0  
Pred : 1, Answer 1  
Pred : 1, Answer 1  
Pred : 0, Answer 0
```

```
Hidden Node Values  
      H1  H2  
[0 0]  0.00 0.06  
[0 1]  0.10 0.98  
[1 0]  0.10 0.98  
[1 1]  0.88 1.00
```

이를 그림으로 나타내면 다음과 같다.

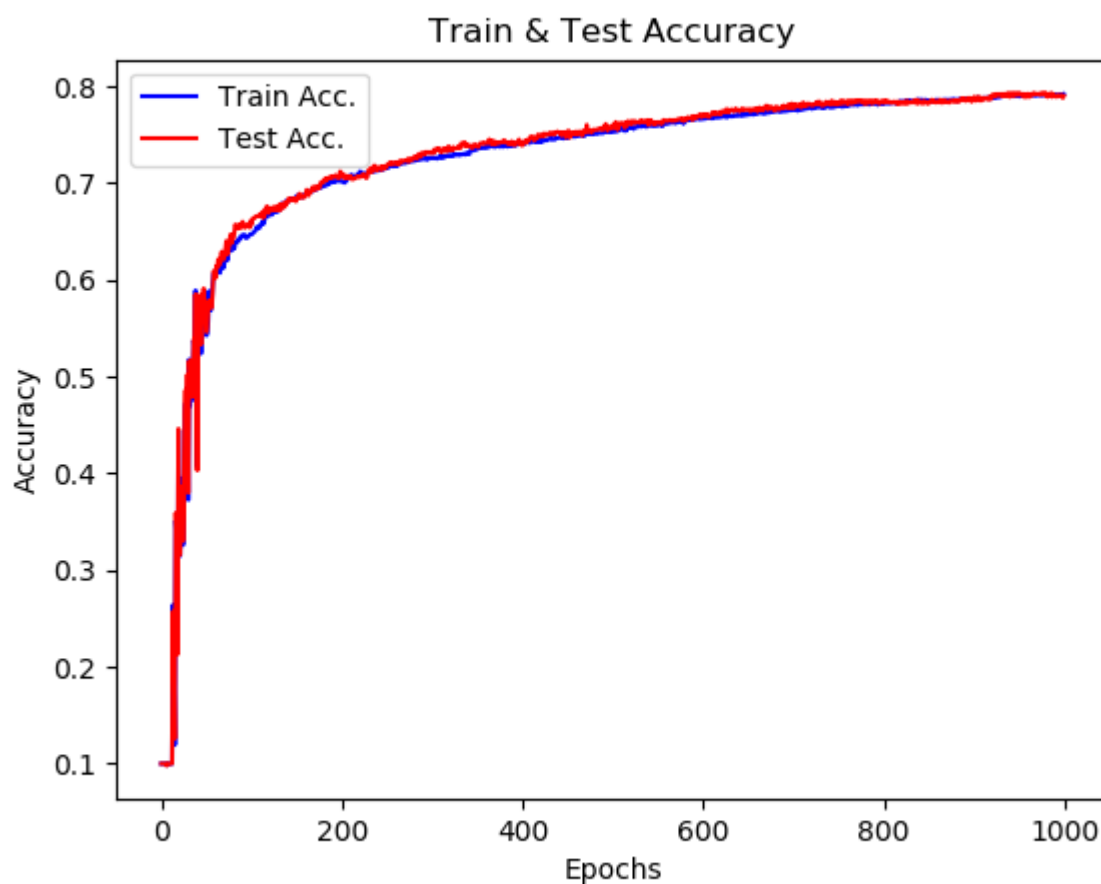


즉, (0,1)과 (1,0)을 같은 hidden node값을 갖게 함으로 인해 다음과 같이 선 하나로 (0,0), (1,1) 과 (0,1), (1,0) 두 영역으로 분리할 수 있도록 구성되었다.

- (b) [MLP-2] Adjust hyperparameters and run 'MLP_2.py' on fashion MNIST to get the best results. Report your best results with the hyperparameters. Show the plot of training and test accuracy according to the number of training epochs with the given code and briefly explain the plot. (batch size = 100)

Hidden 1	Hidden 2	# of epochs	Learning rate	Acc.
100	100	1000	0.0001	0.790

Figure 1



x=5.4381 y=0.69661

초반 약 100 epoch정도에서는 조금 들쭉날쭉한 경향을 보이기는 하지만, accuracy가 급격하게 잘 증가하는 모습을 볼 수 있고, 이후 epoch이 증가함에 따라 Train과 Test accuracy가 증가하지만, 그 속도는 현저하게 낮아진다. 즉, 초반에는 optimal을 찾기 위해 많이 움직인다는 것을 알 수 있고, optimal을 어느 정도 찾은 후에는 안정되게 optimal을 향해 천천히 수렴한다는 것을 알 수 있다. 이는 SGD 외의 다른 Gradient Descent 방법을 사용한다면 더욱 효과적으로 수렴할 것이라 예상한다.

주목할 만한 점은 training과 test accuracy가 모두 증가한다는 것을 알 수 있고, 이는 overfitting등의 문제가 발생하지 않았다는 점을 보여준다.